



ifis

Institut für Informationssysteme
Technische Universität Braunschweig

Data Warehousing & Mining Techniques

Wolf-Tilo Balke

Muhammad Usman

Institut für Informationssysteme
Technische Universität Braunschweig
<http://www.ifis.cs.tu-bs.de>



- Last week:
 - Logical Model: Cubes, Dimensions, Hierarchies, Classification Levels
 - Physical Level
 - Relational: **Star-, Snowflake-schema**
 - Multidimensional (array based storage): **linearization**, problems e.g., **order of dimensions**, dense and sparse cubes
- This week:
 - Indexes



4. Indexes

4. Indexes

4.1 Tree based indexes

4.2 Bitmap indexes





4.0 Indexes

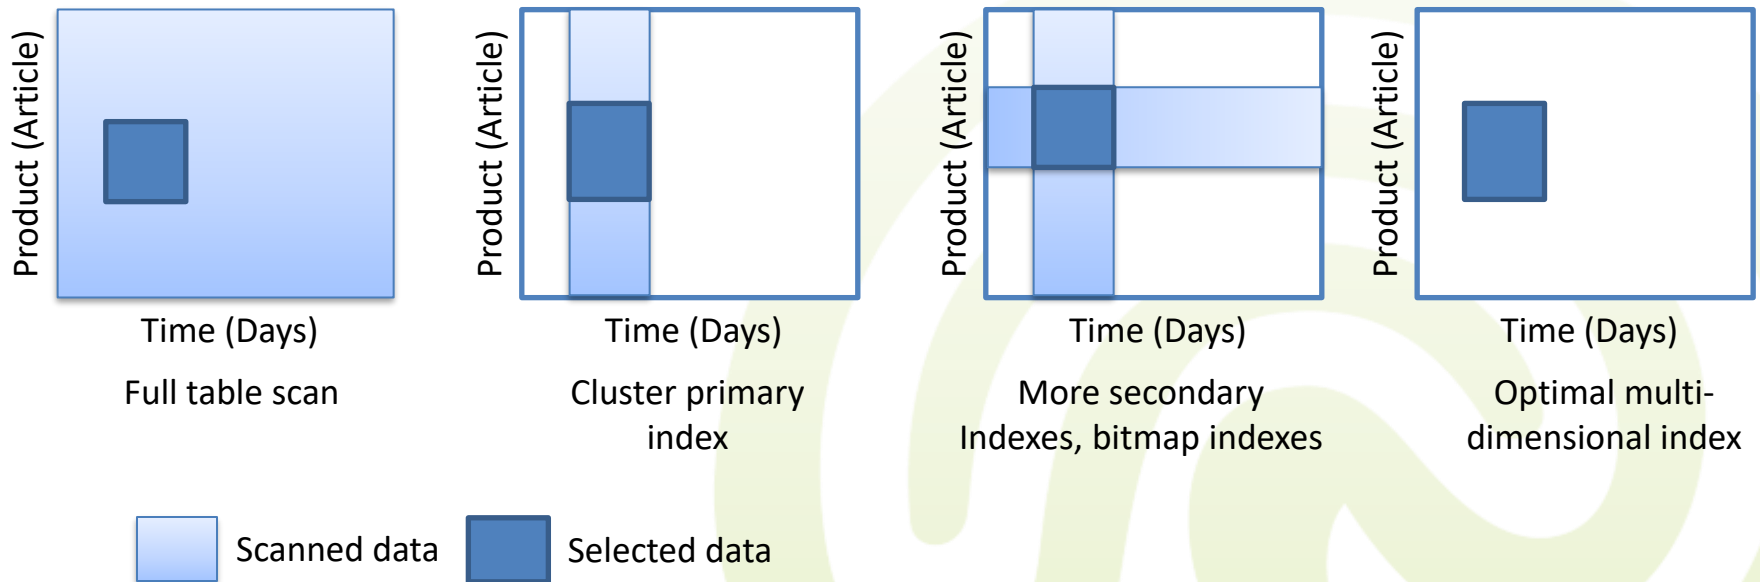
- Why index?
 - Consider a 100 GB table; at 100 MB/s read speed we need 17 minutes for a full table scan
 - Query for the number of Bosch S500 washing machines sold in Germany last month
 - Applying restrictions (product, location) the selectivity would be strongly reduced
 - If we have 30 locations, 10000 products and 24 months in the DW, the selectivity is
$$1/30 * 1/10000 * 1/24 = 0,00000014$$
 - So...**we read 100 GB for 1,4KB of data**





4.0 Indexes


- **Reduce the size of read pages to a minimum with indexes**





4.1 Tree Based Indexes


- In the beginning...there were **B-Trees**
 - Data structures for storing sorted data with amortized run times for insertion and deletion
 - Basic structure of a node





4.1 Tree Structures


- Search in database systems
 - **B-tree** structures allow exact search with logarithmic costs





4.1 Tree Structures


- **Search in DWs**
 - The data is **multidimensional**, B-trees however, support only one-dimensional search
- Are there any possibilities to extend tree functionality for multidimensional data?





4.1 Tree Structures

- **The basic idea** of multidimensional trees
 - Describe the sets of points through **geometric regions**, which contain (clusters of) data points
 - Only clusters are considered for the actual search and not every individual point
 - Clusters can contain each other, resulting in a **hierarchical structure**





4.1 Tree Structures

- Differentiating criteria for tree structures:
 - **Cluster construction:**
 - Either completely fragmenting the space
 - Or grouping data locally
 - **Cluster overlap:**
 - Overlapping or
 - Disjoint
 - **Balance:**
 - Balanced or
 - Unbalanced






4.1 Tree Structures

– Object storage:

- Objects can be in leaves and nodes, or
- Objects are only in the leaves

– Geometry:


- Hyper-sphere,
- Hyper-cube,
- ...





4.1 R-Trees

- The **R-tree** (Guttman, 1984) is a multi-dimensional extension of the classical B-trees
 - Frequently used for low-dimensional applications (up to about 10 dimensions), such as geographic information systems
- More scalable versions: R^+ -Trees, R^* -Trees and X-Trees
 - each up to 20 dimensions for uniformly distributed data




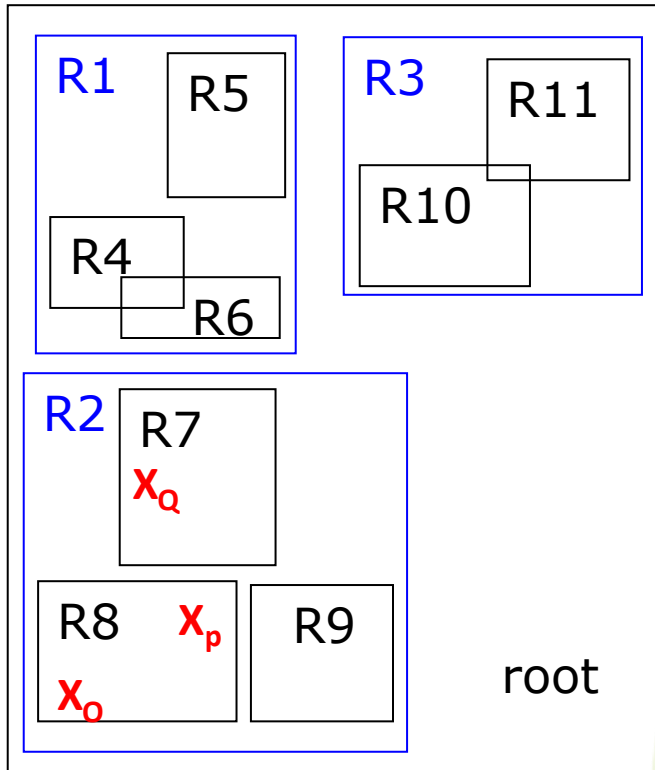


4.1 R-Tree Structure

- **Dynamic Index Structure**
(insert, update and delete are possible)
- **Data structure**
 - **Data pages** are leaf nodes and store clustered point data and data objects
 - **Directory pages** are the internal nodes and store directory entries
 - Multidimensional data are structured with the help of **Minimum Bounding Rectangles (MBRs)**



4.1 R-Tree Example





4.1 R-Tree Characteristics

- **Local grouping** for clustering
- **Overlapping** clusters
 - the more clusters overlap the **less efficient** the index
- **Height-balanced** tree structure
 - therefore all the children of a node in the tree have about the same number of successors
- Objects are **stored** only in the leaves
 - Internal nodes are used for navigation
- MBRs are used as **geometry**



4.1 R-Tree Properties

- The root has **at least two children**
- Each internal node has between m and M children
- M and $m \geq M / 2$ are pre-defined parameters
- All the leaves in the tree are on the **same level**
- All leaves have between m and M index records
- Internal nodes: $(I, \text{child-pointer})$ where I is the **smallest bounding rectangle** that contains the rectangles of the child nodes
- Leaf nodes: $(I, \text{tuple-id})$ I is the smallest bounding rectangle that contains the data object (with ID tuple-id)



4.1 Operations on R-Trees

- The essential operations for the use and management of an R-tree are
 - Search
 - Insert
 - Updates
 - Delete
 - Splitting






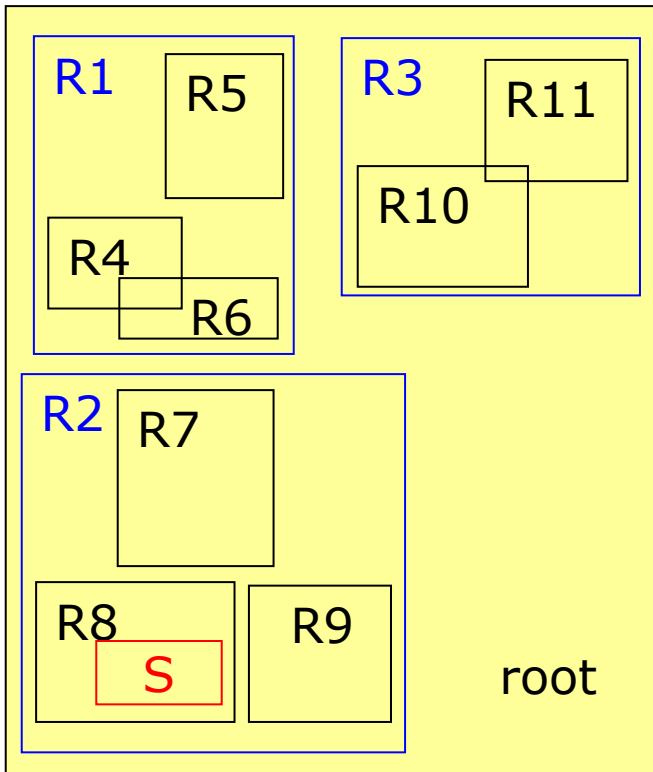
4.1 Searching in R-Trees

- The tree is searched **recursively** from the root to the leaves
 - One path is selected
 - If the requested record has not been found in that sub-tree, the next path is traversed
- The path selection is arbitrary





4.1 Example



Check all the objects
in node R8

- Check only 7 nodes instead of 12



4.1 Searching in R-Trees

- **No guarantee** for good performance
- In the worst case, all paths must be traversed (due to overlaps of the MBRs)
- Search algorithms try to exclude as many irrelevant regions as possible (“pruning”)






4.1 Search Algorithm

- All the index entries which intersect with the **search rectangle (query)** are traversed
 - **The search in internal nodes**
 - Check each MBR for intersection with query
 - For all intersecting MBRs continue the search with their children
 - **The search in leaf nodes**
 - Check all the data points to determine whether they intersect the query
 - Take all correct objects into the result set



4.1 Insert

- Choose the **best leaf** page for inserting the data
 - The best leaf is the leaf that needs the **smallest volume growth** to include the new object
- Why smallest volume growth?
 - Enlarging R1 produces a large portion of unoccupied space in R1 (dead space)
 - Since R1 occupies now large portions of space, the probability of a query intersecting with R1 is bigger, but the probability of hitting real data is low





4.1 Heuristics

- An object is always inserted into nodes, where it produces the smallest increase in volume
 - If it falls in the **interior** of some MBR, no enlargement is needed
 - If there are several possible nodes, select the one with the **smallest overall** volume





4.1 Insert

- After the leaf is chosen, the object is inserted whenever there is still space (remember: number of objects in each node $\leq M$)
 - Otherwise it is considered a case for **overflow handling** and the leaf node is divided
 - The interval (the bounding rectangle) of the parent node must be adapted to the new object
 - Divisions can cascade, if the parent was also full
 - If the root is reached by division, then create a new root whose children are the two split nodes of the old root




4.1 Insert with Overflow




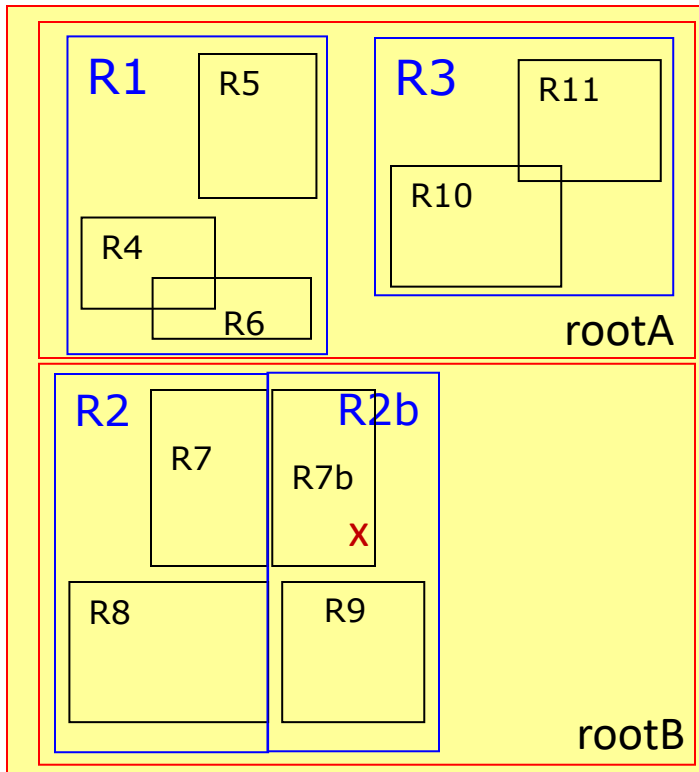
$M = 3$

R7 is full but is chosen
as leaf for insertion
R2 is also full






4.1 Insert with Overflow






4.1 Splitting a Node

- The goal in splitting is that it should rarely be needed to **traverse both resulting nodes** on subsequent searches
 - Therefore use **small** MBRs. This leads to minimal **overlapping** with other MBRs



Bad split



Better Split



4.1 Overflow Problem

- Deciding on how exactly to perform the splits is **not trivial**
 - All objects of the old MBR can be divided in different ways on two new MBRs
 - The volume of both resulting MBRs should remain as small as possible
 - The naive approach of checking all splits and calculate the resulting volumes is not possible
- Two approaches
 - With **quadratic cost**
 - With **linear cost**



4.1 Overflow Problem

- Procedure with **quadratic cost**
 - For each 2 objects compute the necessary MBR and choose the pair with the largest MBR
 - Since these two objects should never occur in some MBR, they will be used as **starting points** for two new MBRs
 - For all other objects compute the **difference** of the necessary **volume increase** to insert them in either one of the starting points
 - Insert the object with the **smallest difference** in the corresponding MBR and compute the MBR again
 - Repeat this procedure for all non-allocated objects

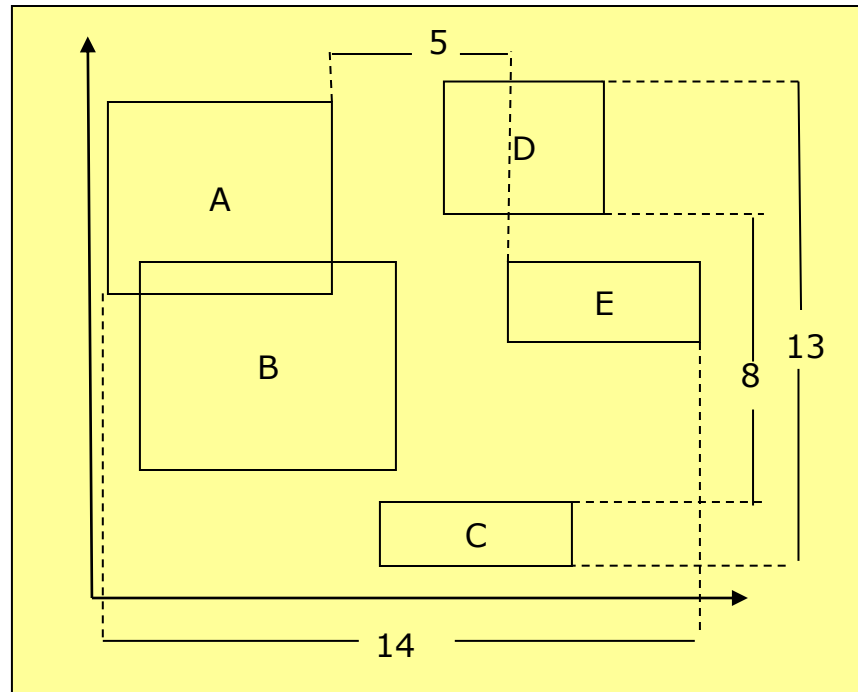


4.1 Overflow Problem

- Procedure with **linear cost**
 - In each dimension:
 - Find the rectangle with the **highest minimum** coordinates, and the rectangle with the **smallest maximum** coordinates
 - Determine the distance between these two coordinates, and normalize it on the size of all the rectangles in this dimension
 - Determine the two starting points of the new MBRs as the two objects with the highest normalized distance



4.1 Example



- x-direction: select A and E, as $d_x = \text{diff}_x / \text{max}_x = 5 / 14$
- y-direction: select C and D, as $d_y = \text{diff}_y / \text{max}_y = 8 / 13$
- Since $d_x < d_y$, C and D are chosen for the split



4.1 Overflow Problem

- Iteratively insert remaining objects in the MBR with the smallest volume growth
- The linear process is a simplification of the quadratic method
- It is usually sufficient providing similar quality of the split (minimal overlap of the resulting MBRs)



4.1 Delete

- **Procedure**




- Search the leaf node with the object to delete
- Delete the object
- The tree is condensed if the resulting node has less than m objects
- When **condensing**, a node is completely erased and the objects of the node which should have remained are **reinserted**
- If the root remains with just one child, the child will become the new root





4.1 Example

- An object from R9 is deleted
(1 object remains in R9, but $m = 2$)
 - Due to few objects R9 is deleted, and R2 is condensed





4.1 Update



- If a record is updated, its surrounding rectangle can change
- The index entry must then be deleted updated and then re-inserted





4.1 Block Access Cost

- The most efficient search in R-trees is performed when the **overlap** and the **dead space** are minimal



Avoiding overlapping is only possible if data points are known in advance





4.1 Improved Versions of R-Trees

- Why may R-trees be inefficient?
 - They **allow overlapping** between neighboring MBRs
- R⁺-Trees (Sellis and others, 1987)
 - **Overlapping of neighboring MBRs is prohibited**
 - This may lead to identical leafs occurring more than once in the tree



4.1 R⁺-Trees



- **Overlaps are not permitted (A and P)**
- Data rectangles are divided and may be present (e.g., G) in **several leaves**



4.1 Performance

- The **main advantage** of R^+ -trees is to improve the search performance
 - Especially for point queries, this may save up to 50% of access time
- **Drawback** is the low occupancy of nodes resulting from many splits
 - R^+ -trees often **degenerate** with the increasing number of changes
 - Actually **scalability** is similar to R-trees



4.1 More Versions

- **R*-trees** and **X-trees** improve the performance of the R^+ -trees (Kriegel and others, 1990/1996)
 - Improved split algorithm in R^* -trees
 - “Extended nodes” in X-trees allow sequential search of larger objects
 - Scalable up to 20 dimensions



4.1 UB-Trees

Detour


- B-Trees are great for indexing uni-dimensional data, but in the DW the data is stored **multidimensional**
- Idea: represent multidimensional data with just one dimension without information loss
 - How? Like in the case of MOLAP, with **linearization**



4.1 UB-Trees

Detour

- Universal B-Tree (UB-tree) are a combination of B⁺-Tree and Z-curve as linearization function
 - Z-curve is used to map multidimensional points to one-dimensional values (Z-values)
 - Z-values are used as keys in B⁺-Tree





4.1 Z-Curve Function

Detour

- **Z-Value address representation**
 - Calculate the z-values such that neighboring data is clustered together
 - Calculated through bit **interleaving** of the coordinates of the tuple
 - In order to localize a value with coordinates one has to perform de-interleaving

For Z-value 51, we have
51 - offset = 50
50 in binary is 110010

Z-value = 110010
Y = 101 = 5 X = 100 = 4

Tuple = 51, x = 4, y = 5

	X							
	0	1	2	3	4	5	6	7
0	1	2	5	6	17	18	21	22
1	3	4	7	8	19	20	23	24
2	9	10	13	14	25	26	29	30
3	11	12	15	16	27	28	31	32
4	33	34	37	38	49	50	53	54
5	35	36	39	40	51	52	55	56
6	41	42	45	46	57	58	61	62
7	43	44	47	48	59	60	63	64



4.1 UB-Trees

Detour

- **Z-Regions**

- The space covered by an interval on the Z-Curve
 - E.g. [1:9], [10, 18], [19, 28]...
- Each Z-Region maps exactly onto **one page on secondary storage**
 - i.e., to one leaf page of the B⁺-Tree
- This allows for very efficient processing of **multidimensional range queries**

1	2	5	6	17	18	21	22
3	4	7	8	19	20	23	24
9	10	13	14	25	26	29	30
11	12	15	16	27	28	31	32
33	34	37	38	49	50	53	54
35	36	39	40	51	52	55	56
41	42	45	46	57	58	61	62
43	44	47	48	59	60	63	64



4.1 UB-Trees

Detour

- Range queries (RQ) in UB-Trees
 - Each query can be specified by 2 coordinates
 - q_a (the upper left corner of the query rectangle)
 - q_b (the lower right corner of the query rectangle)
 - RQ-algorithm
 - I. Starts with q_a and calculates its Z-Region
 - I. Z-Region of q_a is [10:18]

1	2	5	6	17	18	21	22
3	4	7	8	19	20	23	24
9	10	13	14	25	26	29	30
11	12	15	16	27	28	31	32
33	34	37	38	49	50	53	54
35	36	39	40	51	52	55	56
41	42	45	46	57	58	61	62
43	44	47	48	59	60	63	64



4.1 UB-Trees

Detour

- Range queries (RQ) in UB-Trees
 2. The corresponding page is loaded and filtered with the query predicate
 - E.g. value 10 has after de-interleaving $x=1$ and $y=2$, which is outside the query rectangle
 3. After q_a all values on the Z-curve are de-interleaved and checked by their coordinates
 - The data is only accessed from the disk
The next jump point on the Z-curve is 27
 4. Repeat steps 2 and 3 until the end-address of the last filtered region is bigger than q_b

1	2	5	6	17	18	21	22
3	4	7	8	19	20	23	24
9	10	13	14	25	26	29	30
11	12	15	16	27	28	31	32
33	34	37	38	49	50	53	54
35	36	39	40	51	52	55	56
41	42	45	46	57	58	61	62
43	44	47	48	59	60	63	64



4.1 UB-Trees

Detour


- The critical part of the algorithm is **calculating the jump point** on the Z-curve which is inside the query rectangle
 - If this takes too long it eliminates the advantage obtained through optimized disk access
 - Calculating the jump point mostly involves
 - Performing **bit operations** and comparisons
 - 3 points: q_a , q_b and the current Z-Value



4.1 UB-Trees

Detour

- Advantages of UB-Trees
 - B⁺-Trees provide for high node filling degree (at least 50%)
 - Logarithmical complexity at **search, insert and delete**
 - The Z-curve provides for good performance for **range queries!**
 - Consecutive values on the Z-curve index similar data
 - Similarity by means of neighborhood





4.2 Bitmap Indexes

- Database indexes that use **bitmaps**
- Bitmaps (bit arrays) are array data structures that store individual bits
- Bitmap indexes are primarily intended for DW
 - Users query data rather than update it
- Bitmap indexes work well for data which has a small number of distinct values
 - E.g. gender data, or dimensions

Identifier	Gender	Bitmaps	
		F	M
1	Female	1	0
2	Male	0	1
3	Male	0	1
4	Unspecified	0	0
5	Female	1	0



4.2 Bitmap Indexes

- Let's assume fact table Sales and dimension Geography with granularity Shops

Nr	Shop
1	Saturn
2	Real
3	P&C

Nr	Shop_ID	Sum
1	1	150
2	2	65
3	3	160
4	2	45
5	1	350
6	2	80

- A bitmap index on the fact table for dimension geography on attribute Shop looks like this

Value	Bitmap
P&C	001000
Real	010101
Saturn	100010



4.2 Bitmap Indexes

- A **bitmap index** for an attribute is:
 - A **collection of bit-vectors**
 - The number of bit-vectors represents **the number of distinct values** of the attribute in the relation
 - The **length of each bit-vector** is called the cardinality of the relation
 - The bit-vector for value 'Saturn' has 1 in position 5, if the 5th record has 'Saturn' in attribute Shop, and it has 0 otherwise
- Records are allocated permanent numbers
- There is a mapping between record numbers and record addresses



4.2 Bitmap Indexes

- Let's assume that...
 - There are n records in the table
 - Attribute A has m distinct values in the table
- The size of a bitmap index on attribute A is $m \cdot n$
- Significant number of 0's if m is big, and of 1's if m is small
 - Opportunity to compress
 - Run Length Encoding (RLE)
 - Gzip (Lempel-Ziv, LZ)
 - Byte-Aligned Bitmap Compression (BBC): variable byte length encoding (Oracle patent)



4.2 Bitmap Indexes

- Handling modification
 - Assume **record numbers are not changed**
- Deletion
 - **Tombstone** replaces deleted record (6 doesn't become 5!)
 - Corresponding bit is set to 0

Nr	Shop	Sum
1	Saturn	150
2	Real	65
3	P&C	160
4	Real	45
5	Saturn	350
6	Real	80

Before

Value	Vector
P&C	001000
Real	010101
Saturn	100010



After

Value	Vector
P&C	001000
Real	010101
Saturn	100000



4.2 Bitmap Indexes

- Inserted record is assigned the next record number
 - A bit of value 0 or 1 is appended to each bit vector
 - If new record contains a new value of the attribute, add one bit-vector
 - E.g., insert new record with REWE as shop

Nr	Shop	Sum
1	Saturn	150
2	Real	65
3	P&C	160
4	Real	45
5	Saturn	350
6	Real	80
7	REWE	23

Before

Value	Vector
P&C	001000
Real	010101
Saturn	100010



After

Value	Vector
P&C	0010000
Real	0101010
Saturn	1000100
REWE	0000001



4.2 Bitmap Indexes

- Performing updates
 - Change the bit corresponding to the old value of the modified record to 0
 - Change the bit corresponding to the new value of the modified record to 1
 - If the new value is a new value of A, then insert a new bit-vector: e.g., replace Shop for record 2 to REWE

Nr	Shop	Sum
1	Saturn	150
2	REWE	65
3	P&C	160
4	Real	45
5	Saturn	350
6	Real	80

Before

Value	Vector
P&C	001000
Real	010101
Saturn	100010



After

Value	Vector
P&C	001000
Real	000101
Saturn	100010
REWE	010000



4.2 Bitmap Indexes

- Performing selects
 - Basic AND, OR bit operations:
 - E.g., select the sums we have spent in either Saturn or P&C

Nr	Shop	Sum
1	Saturn	150
2	Real	65
3	P&C	160
4	Real	45
5	Saturn	350
6	Real	80

Saturn	OR	P&C	=	Result
1		0		1
0		0		0
0		1		1
0		0		0
1		0		1
0		0		0

Value	Vector
P&C	001000
Real	010101
Saturn	100010

- Bitmap indexes should be used when selectivity is **high**



4.2 Bitmap Indexes

- Advantages
 - Operations are efficient and easy to implement (directly supported by hardware)
- Disadvantages
 - For each new value of an attribute a new bitmap-vector is introduced
 - If we bitmap index an attribute like birthday (only day) we have 365 vectors: $365/8 \text{ bits} \approx 46 \text{ Bytes}$ for a record, just for that
 - Solution to such problems is **multi-component** bitmaps
 - Not fit for range queries where many bitmap vectors have to be read
 - Solution: **range-encoded** bitmap indexes



4.2 Bitmap Indexes

- **Multi-component bitmap indexes**
 - Encoding using a different numeration system
 - E.g., for the month attribute, between 0 and 11 values can be encoded as $x = 4 * y + z$, where $0 \leq y \leq 2$, and $0 \leq z \leq 3$, called $\langle 3,4 \rangle$ basis encoding

Month	Dec	Nov	Oct	Sep	Aug	Jul	Jun	Mai	Apr	Mar	Feb	Jan
M	A_{11}	A_{10}	A_9	A_8	A_7	A_6	A_5	A_4	A_3	A_2	A_1	A_0
5	0	0	0	0	0	0	1	0	0	0	0	0

- $5 = 4 * 1 + 1$

X	Y			Z			
M	$A_{2,1}$	$A_{1,1}$	$A_{0,1}$	$A_{3,0}$	$A_{2,0}$	$A_{1,0}$	$A_{0,0}$
5	0	1	0	0	0	1	0



4.2 Multi-Component Bitmap Indexes

- **Advantage of multi-component bitmap indexes**
 - If we have 100 (0..99) different days to index we can use a multi-component bitmap index with basis of $\langle 10, 10 \rangle$
 - The storage is reduced from 100 to 20 bitmap-vectors (10 for y and 10 for z)
 - The read-access for a point (1 day out of 100) query needs however 2 read operations instead of just 1



4.2 Bitmap Indexes

- **Range-encoded bitmap indexes:** Persons born between March and August
 - For normal encoded bitmap indexes read 6 vectors

	Dec	Nov	Oct	Sep	Aug	Jul	Jun	Mai	Apr	Mar	Feb	Jan
Person	A_{11}	A_{10}	A_9	A_8	A_7	A_6	A_5	A_4	A_3	A_2	A_1	A_0
1	0	0	0	0	0	0	1	0	0	0	0	0
2	0	0	0	0	0	0	0	0	1	0	0	0
3	0	0	0	0	0	0	0	0	0	0	0	1
4	0	0	0	0	0	0	0	0	1	0	0	0
5	1	0	0	0	0	0	0	0	0	0	0	0

- Idea: set the bits of all bitmap vectors to 1 if they are higher or equal to the given value



4.2 Bitmap Indexes

- **Query:** Persons born between March and August
 - So persons which didn't exist in February, but existed in August!
 - **Just 2 vectors** read: $((\text{NOT } A_1) \text{ AND } A_7)$

	Dec	Nov	Oct	Sep	Aug	Jul	Jun	Mai	Apr	Mar	Feb	Jan
Person	A_{11}	A_{10}	A_9	A_8	A_7	A_6	A_5	A_4	A_3	A_2	A_1	A_0
1	1	1	1	1	1	1	1	0	0	0	0	0
2	1	1	1	1	1	1	1	1	1	0	0	0
3	1	1	1	1	1	1	1	1	1	1	1	1
4	1	1	1	1	1	1	1	1	1	0	0	0
5	1	0	0	0	0	0	0	0	0	0	0	0



4.2 Range-encoded Bitmap Indexes

- If the query is limited only on one side, (e.g., persons born in or after March), **Ivector** is enough (NOT A_1)
- For point queries, **2 vector reads** are however necessary!
 - E.g., persons born in March: $((\text{NOT } A_1) \text{ AND } A_2)$



4.2 Advantages

- Bitmap indexes are great for indexing the dimensions
 - Fully indexing tables with a traditional trees can be expensive - the indexes can be several times larger than the data
 - Bitmap indexes are typically only a fraction of the size of the indexed data in the table.
- They...
 - reduced response time for large classes of ad hoc queries
 - bring dramatic performance gains even on hardware with a relatively small number of CPUs or a small amount of memory



- B-Trees are not fit for multidimensional data
- R-Trees
 - MBR as geometry to build multidimensional indexes
 - Operations: select, insert, overflow problem, node splitting , delete
 - Inefficient because they **allow overlapping** between neighboring MBRs
 - R⁺-trees - **improve** the search performance



- UB-Trees
 - Reduce multidimensional data to one dimension in order to use B-Tree indexes
 - Z-Regions, Z-Curve, use the advantage of bit operations to make optimal jumps
- Bitmap indexes
 - Great for indexing tables with set-like attributes e.g., Gender: Male/Female
 - Operations are efficient and easy to implement (directly supported by hardware)
 - **Multi component** reduce the storage while **range encoded** allow for fast range queries



Next lecture

- Optimization
 - Partitioning
 - Joins
 - Materialized Views

