



**ifis**

Institut für Informationssysteme  
Technische Universität Braunschweig

# Data Warehousing & Mining Techniques

**Wolf-Tilo Balke**

**Muhammad Usman**

Institut für Informationssysteme  
Technische Universität Braunschweig  
<http://www.ifis.cs.tu-bs.de>



- Last Week: DW-Project
- ETL Process
- This Lecture: Real-Time DW



# 8. Real-Time DW

## 8. Real-Time Data Warehouses

### 8.1 Real-Time Requirements

### 8.2 Enabling Real-Time ETL

### 8.3 OLAP and the Changing Data

### 8.4 OLTP & OLAP - Hybrid Solutions



A photograph of a Deutsche Bahn (DB) departure board. The board is divided into columns for 'Départ', 'Gleis' (Platform), and 'Hinweis' (Note). The 'Départ' column shows train destinations like 'S', 'E N', 'M', and 'G'. The 'Gleis' column shows platform numbers like 12, 2, 12, 12, 13, 12, and 11. The 'Hinweis' column contains messages such as 'Bitte Ansage beachten' and 'über 60 Min später'.

| Départ | Gleis | Hinweis               |
|--------|-------|-----------------------|
| S      | 12    | Bitte Ansage beachten |
|        | 2     | über 60 Min später    |
|        | 12    | über 60 Min später    |
|        | 12    | über 60 Min später    |
| E N X  | 13    | über 60 Min später    |
| M      | 12    |                       |
| G      | 11    |                       |





## 8.1 Real-Time DW - Why?


- **Customer business scenario:** a utility company owns plants generating energy
- Existing DW supports planning by recommending:
  - The production capacity
  - The reserve capacity
  - When to buy supplemental energy, as needed





## 8.1 Real-Time DW - Why?

- Each day is pre-planned on **historical behavior**
  - Peak demand periods are somewhat predictable
- Good planning is important because:
  - Expensive to have unused capacity!
  - Cheaper to buy energy ahead!
- **Planning on last week's average is not enough**





## 8.1 Real-Time DW - Why?

- Getting more **in-time** accuracy enhances operational business
  - Compare **today's** plant output and customer consumption volumes to yesterday's or last week's average
  - Know when to purchase additional options or supplies
- Customer Target: have the actual data from the operational environment available for analytics within a **5 minute lag**
- **Real-time**  $\neq$  **fast**
  - Real time DW has the capability to **enforce** time constraints



## 8.1 Real-Time DW - Why?

- The most difficult part for complying with the 5 minutes time constraint is the **ETL process**
  - ETL tools usually operate in batch mode on a certain schedule nightly, weekly or monthly
  - ETL typically involves **downtime** for the DW during the loading step (usually happens over night)
  - Data is extracted into flat files in the staging area **outside the DBMS**, where it is not available for querying
  - ETL may take **hours** to finish
- The ETL process needs to be re-designed to meet real-time constraints





## 8.2 Real-Time ETL

- Solutions enabling **real-time ETL**
  - Microbatch
  - Direct Trickle-feed
  - Trickle & Flip
  - External Real-time Data Cache (ERDC)







## 8.2 Microbatches

- **Microbatches** relies on the classical ETL-batch solution to provide **near-real-time ETL**
  - Reduce the time interval between consecutive loads (typical intervals are 3-4 hours)
  - Transformation tasks have to be fully automatized (no human intervention)
  - However, the time interval can't be reduced to minutes
  - The interval depends on:
    - The **data volume**
    - The **operational system and OLAP load** during the microbatch process



## 8.2 Direct Trickle-Feed

- **Direct Trickle-Feed** (also called continuous feed)
  - **Continuously** feed the DW with new data from the operational environment
    - Directly propagate changes on data from the OD store as inserts into the DW fact table
  - But **constantly updating** the same tables being queried by a **reporting or OLAP** tool can cause the DW's query **performance** to degrade – **query contention**
    - Under moderate to heavy usage from either OLAP queries or the incoming data, most RDMS will block the incoming data – the **Walmart scenario**



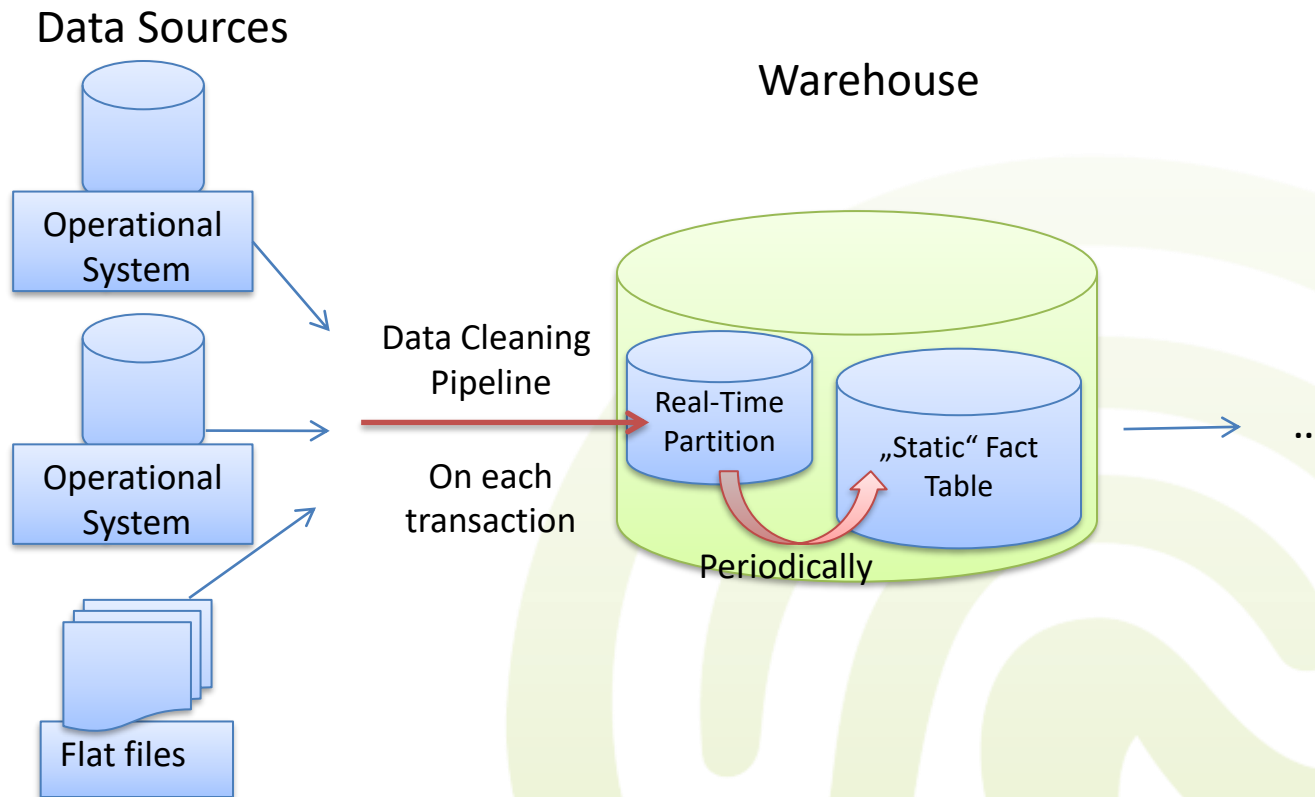
## 8.2 Direct Trickle-Feed

- Solution: **separate real-time fact partition**
- Idea is using a separate table subject to special rules for update and query
  - Contains all the activity that occurs since the last update of the “static” (updated each night) fact table
  - Linked as seamlessly as possible to the content of the “static” fact table
  - Indexed only lightly to account for incoming data
  - Support highly responsive queries



## 8.2 Direct Trickle-Feed


- Real-time partition architecture





## 8.2 Real-Time Partition in Practice

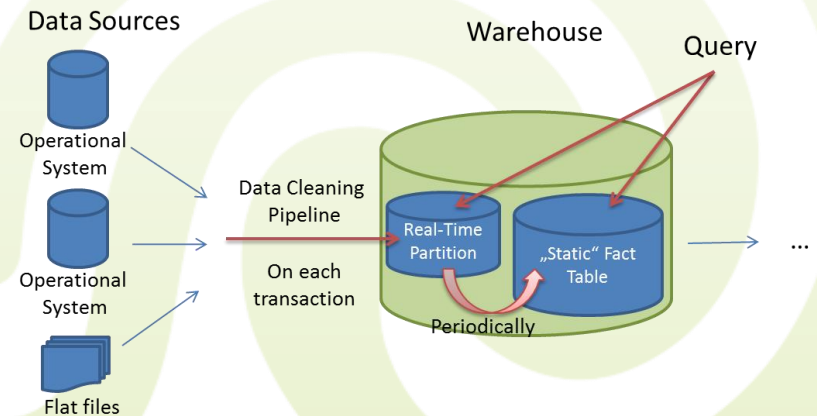
- The Amazon DW:
  - An estimate of 55 GB sales transactions a year
    - For a day, this would mean about 150 MB of raw data
  - The “static” fact table just for the sales cube, for the last 3 years, would be about 150 GB
    - For fast OLAP querying it will be heavily indexed and supported by aggregates
  - The small 150 MB real-time partition can be pinned in memory for rapid loading and querying





## 8.2 Real-Time Partition in Practice

- The query: How did the sales for today evolve until now (Tuesday the 6<sup>th</sup> of December, at 10:30) compared to the last 3 Tuesdays?
  - The **query processor** detects that **fresh data** is required by the query
  - The query processor sends the same query to both the “static” fact table and to the real-time partition





## 8.2 Real-Time Partition in Practice

- Data from the real-time partition is aggregated to the necessary granularity in concordance with the classification schema
- The system benefits from the **indexes** on the “static” fact table, and from the small **size** of the in-memory real-time partition
- As a downside, this approach may still suffer from **query contention**, if daily data volume is high





## 8.2 Real-Time ETL

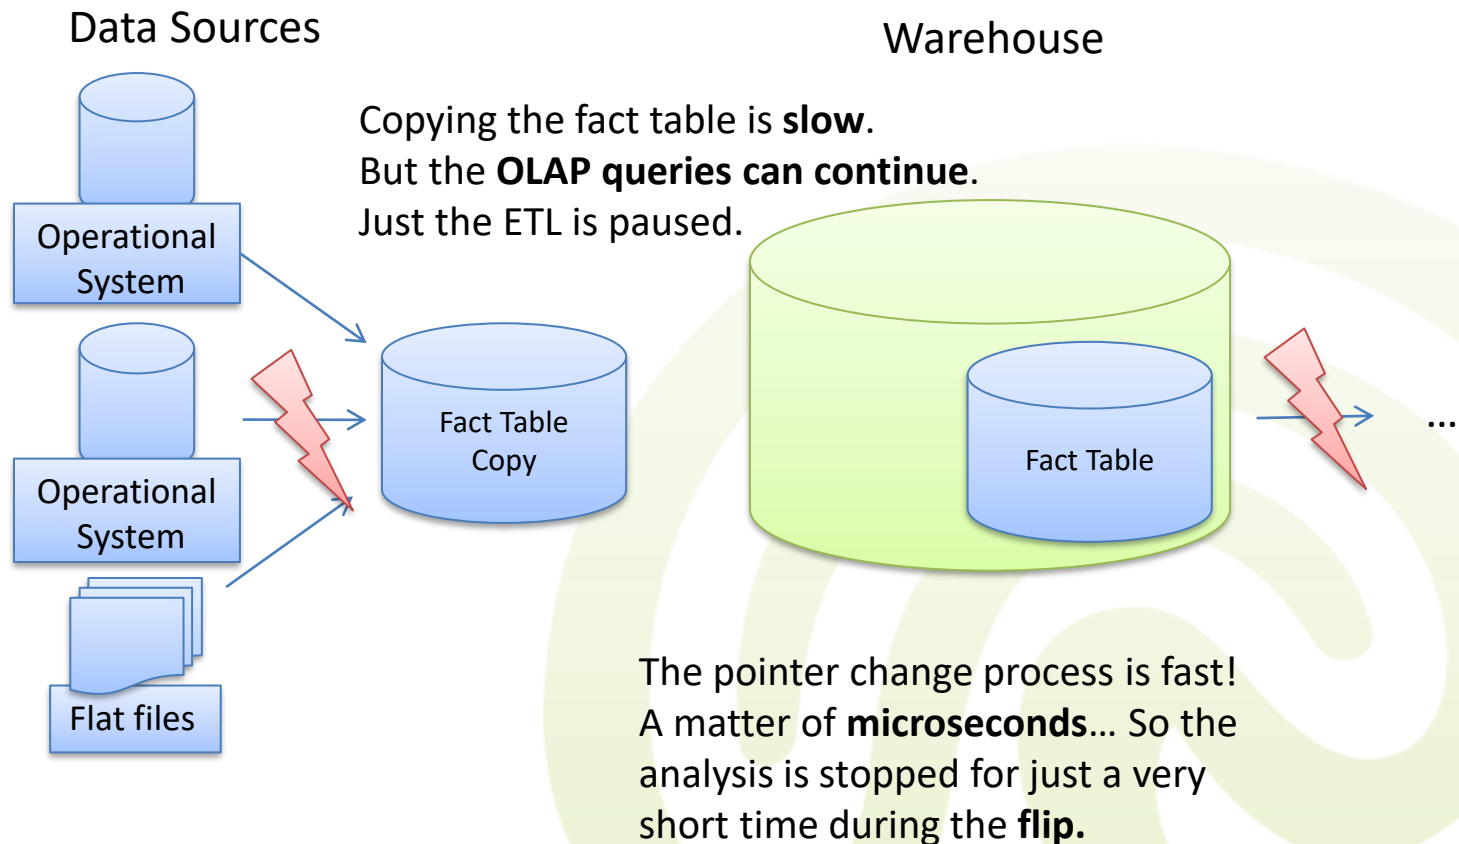
- **Trickle & Flip**

- Instead of continuously feeding the data into the DW, the data is fed into the staging tables in the **exact same format as the target fact tables**
  - This way the staging tables are copies of the fact tables, with up-to-date data
- Periodically, the trickle is halted, the staging tables are copied, and renamed to the active fact table names (just a pointer switch operation)
  - The active fact table is deleted, and the process begins anew



## 8.2 Trickle & Flip

- Architecture:





## 8.2 Trickle & Flip

- The main problem is **scalability**
  - The staging tables must contain the **complete fact table**
  - If the fact table is too large the copying process is slow
  - This method is only useful if the frequency of refreshs corresponds to the time it takes to perform the flip
- **No query contention** between complex OLAP queries and loading inserts
  - During flipping queries are paused



## 8.2 ERDC

- **External Real-Time Data Cache (ERDC)**
  - Leave the DW largely as-is
  - Store the incoming real-time data outside the traditional DW in order to completely avoid performance problems
  - The data cache can simply be a dedicated database server for loading storing and processing real-time data





## 8.2 ERDC

- Two approaches querying the real-time data
  - **Separate & isolate** in the real-time data cache
    - All queries involving real-time data are redirected towards the ERDC – similar procedure as the direct trickle-feed
  - **Just-in-time information merge** from external data cache (JIM)
    - The real-time data required to answer any particular query is seamlessly imaged to the regular DW on a temporary basis – similar to trickle & flip



## 8.2 Separate & Isolate

- All the real-time data loading and querying is routed to the external database
  - No scalability problem for the DW itself
  - The assumption is that the real-time daily data volume is small
  - Data is regularly (e.g., every night) loaded from the ERDC into the DW
- With the real-time data external to the DW it is not possible for a single report to **join real-time and historical** information



## 8.2 Just-in-Time Merge


- An extension of the Separate & Isolate method
- Workflow at query time
  - Queries are pre-processed to determine which **real-time data** parts are **required** (attributes & rows)
  - A snapshot of these required parts is taken and loaded as temporary tables in the DW
  - Once the tables are loaded in the DW, the **query is re-written** to also include the data from the temporary tables





## 8.2 Just-in-Time Merge

- Architecture:





## 8.2 Just-in-Time Merge

- Advantages of JIM
  - Less scalability problems as the real-time data is brought into the DW only on request
  - **Query contention** is not a problem as the data in the temporary tables are snapshots and **do not change while queried**
- Problem: the query processor is very complex



## 8.2 Real-Time ETL

- **Choose** one of the real-time ETL enabling methods based on your needs!
  - Microbatch offers the easy/cheap way out, when near real-time ETL satisfies the requirements (3-4 hours)
  - Direct Trickle-feed with real-time partitions offers real-time performance
    - It works fine for normal daily data loads
    - But may suffer from **query contention** for complex OLAP queries and large number of users





## 8.2 Real-Time ETL

- Trickle & Flip offers near real-time performance
  - The size of the delay is dictated by the size of the fact table – more precisely by the overhead of making a copy of the fact table
- External Real-time Data Cache (ERDC) with Just-in-Time Merge offers true real-time performance while allowing for
  - Rapidly changing data (10-1000 transactions/sec from the ODS)
  - Concurrent access for hundreds of users
  - Conjunction of real-time and historical data
- ERDC with Just-in-Time Merging sounds great...but **it's not supported by any product yet!**





## 8.3 OLAP on Changing Data

- OLAP queries were designed to operate on top of **unchanging, static, historical data**
  - No precautions are taken to ensure **consistency**
    - Data changes concurrent to execution may negatively influence the results of OLAP queries
  - This leads to **inconsistent results** and **confusing reports**





## 8.3 Example: OLAP on Changing Data

- For a *sales by product category* report an OLAP tool issues the following multi-pass SQL statement

```
0:00 create table TEMP1 (  
    Category_Id LONG, DOLLARSALES DOUBLE)  
  
0:01 insert into TEMP1  
    select      a11.[Category_Id] AS Category_Id,  
               sum(a11.[Tot_Dollar_Sales]) as DOLLARSALES  
    from        [YR_CATEGORY_SLS] a11  
    group by    a11.[Category_Id]  
  
0:05 create table TEMP2 ( ALLPRODUCTSD DOUBLE)  
  
0:06 insert into TEMP2  
    select sum(a11.[Tot_Dollar_Sales]) as  
    ALLPRODUCTSD  
    from    [YR_CATEGORY_SLS] a11  
  
0:08 select distinct pa1.[Category_Id] AS Category_Id,  
    a11.[Category_Desc] AS Category_Desc,  
    pa1.[DOLLARSALES] as DOLLARSALES,  
    (pa1.[DOLLARSALES] / pa2.[ALLPRODUCTSD])  
    as DOLLARSALESC  
    from      [TEMP1] pa1,  
             [TEMP2] pa2,  
             [LU_CATEGORY] a11  
    where     pa1.[Category_Id] =  
    a11.[Category_Id]  
  
0:09 drop table TEMP1  
0:10 drop table TEMP2
```

| Category      | Metrics Dollar Sales | Dollar Sales Contribution to all Products Abs. |
|---------------|----------------------|--|
| Electronics   | \$39,915.00          | 19.7%  |
| Food          | \$10,938.00          | 5.4%   |
| Gifts         | \$36,362.00          | 18.0%  |
| Health&Beauty | \$11,707.00          | 5.8%   |
| Household     | \$88,774.00          | 43.8%  |
| Kid's Corner  | \$5,502.00           | 2.7%   |
| Travel        | \$9,289.00           | 4.6%   |
| Total         | \$202,487.00         | 100.0%   |



## 8.3 Example: OLAP on Changing Data

- Lets assume few larger sales are inserted in the DW between seconds 1 and 6
  - The sales by category from TEMP1 won't match the sum from TEMP2
  - The report will then show inconsistent values

| Category      | Metrics Dollar Sales | Dollar Sales Contribution to all Products Abs. |
|---------------|----------------------|--|
| Electronics   | \$39,915.00          | 19.2%  |
| Food          | \$10,938.00          | 5.3%   |
| Gifts         | \$36,362.00          | 17.5%  |
| Health&Beauty | \$11,707.00          | 5.6%   |
| Household     | \$88,774.00          | 42.8%  |
| Kid's Corner  | \$5,502.00           | 2.7%   |
| Travel        | \$9,289.00           | 4.5%   |
| Total         | \$207,487.00         | 97.6%  |

```
0:00 create table TEMP1 (  
    Category_Id LONG, DOLLARSALES DOUBLE)  
  
0:01 insert into TEMP1  
select    a11.[Category_Id] AS Category_Id,  
          sum(a11.[Tot_Dollar_Sales]) as DOLLARSALES  
from      [YR_CATEGORY_SLS] a11  
group by  a11.[Category_Id]  
  
0:05 create table TEMP2 ( ALLPRODUCTSD DOUBLE)  
  
0:06 insert into TEMP2  
select sum(a11.[Tot_Dollar_Sales]) as  
ALLPRODUCTSD  
from    [YR_CATEGORY_SLS] a11  
  
0:08 select distinct pa1.[Category_Id] AS Category_Id,  
a11.[Category_Desc] AS Category_Desc,  
pa1.[DOLLARSALES] as DOLLARSALES,  
(pa1.[DOLLARSALES] / pa2.[ALLPRODUCTSD])  
as DOLLARSALESC  
from      [TEMP1] pa1,  
          [TEMP2] pa2,  
          [LU_CATEGORY] a11  
where     pa1.[Category_Id] =  
a11.[Category_Id]  
  
0:09 drop table TEMP1  
0:10 drop table TEMP2
```

Sales of  
\$5000  
for different  
products





## 8.3 OLAP on Changing Data

- Is such a scenario possible?
  - OLAP queries may have response times ranging from seconds up to hours!
  - It is not uncommon for OLAP reports to consist of 10-50 passes!
- Such a scenario is not only possible but **most probable** and a **problem for real-time DW**
  - True real-time solutions like the Direct Trickle-feed with Real-Time Partition are affected by this problem



## 8.3 OLAP on Changing Data

- Real-time ETL solutions not affected by the inconsistency problem
  - Near real-time solutions like Microbatch and Trickle & Flip
    - Both work on “fresh” **copies** of data which are **not refreshed during OLAP queries**
  - The real-time solution ERDC Just-in-Time solution is also not affected
    - Because data from the external cache is **shadowed on demand** in the DW



## 8.4 Back to Basics

*Detour*

- OLTP and OLAP: How different are they?
  - Both involve a large number of selects
  - More range lookups for the OLAP and fewer inserts updates and deletes
  - But the OLAP queries span over **larger amounts of data** and **take longer** to resolve



## 8.4 Hybrid Solutions

*Detour*

- The approaches we have discussed until now are all **workarounds** for the **locking** constraints combined with **poor hardware performance**
  - OLAP queries take too long to execute
  - OLTP is locked out while performing OLAP
- Simple concept: Shadow the OLTP data in real-time, for OLAP to run on those shadows
  - Avoid **inconsistencies**
  - Limited by **yesterday's hardware**
- **Can we build a hybrid OLTP & OLAP solution with better hardware?**



## 8.4 Where is the Bottleneck?

# Detour

- **The bottleneck:** Slowest components in the computer are usually the **hard drives**
  - Most of the time, the CPU waits for data to be read from the HDD into the random access memory (RAM)
- Windows boots up to 3 times faster when installed on an SSD drive than on a HDD





## 8.4 Storage Media

*Detour*

- Data is stored on a **storage media**. Media highly differ in terms of
  - Random **Access** Speed
  - Random/ Sequential **Read/Write** speed
  - **Capacity**
  - **Cost** per Capacity





## 8.4 Storage Media

*Detour*

- **Random Access Time:** Average time to access a random piece of data at a known media position
  - Usually measured in ms or ns
  - Within some media, access time can vary depending on the position (e.g. hard disks)
- **Transfer Rate:** Average amount of consecutive data which can be transferred per time unit
  - Usually measured in KB/sec, MB/sec, GB/sec,...





## 8.4 Storage Media

*Detour*

- **Volatile:** Memory needs constant power to keep data
  - **Dynamic:** Dynamic volatile memory needs to be “refreshed” regularly to keep data
  - **Static:** No refresh necessary
- **Access Modes**
  - **Random Access:** Any piece of data can be accessed in approximately the same time
  - **Sequential Access:** Data can only be accessed in sequential order



## 8.4 Storage Media

*Detour*

- Media characteristics result in a **storage hierarchy**
- DBMS optimize data distribution among the storage levels
  - **Primary Storage:** Fast, limited capacity, high price, usually volatile electronic storage
    - Frequently used data / current work data
  - **Secondary Storage:** Slower, large capacity, lower price
    - Main stored data
  - **Tertiary Storage:** Even slower, huge capacity, even lower price, usually offline
    - Backup and long term storage of not frequently used data



## 8.4 Storage Media

*Detour*

| Type | Media  | Size   | Random Acc. Speed | Transfer Speed | Characteristics    |
|------|--|--------|-------------------|----------------|--------------------|
| Pri  | L1-Processor Cache<br>(Intel QX9000 )        | 32 KiB | 0.0008 ms         | 6200 MB/sec    | Vol, Stat, RA,OL   |
| Pri  | DDR3-Ram<br>(Corsair 1600C7DHX)              | 16 GiB | 0.004 ms          | 8000 MB/sec    | Vol, Dyn, Ra, OL   |
| Sec  | Harddrive SSD<br>(OCZ Vertex2)               | 22 TB  | < 1 ms            | 285 MB/sec     | Stat, RA, OL       |
| Sec  | Harddrive Magnetic<br>(Seagate ST32000641AS) | 3 TB   | 8.5 ms            | 138 MB/sec     | Stat, RA, OL       |
| Ter  | DVD+R<br>(Verbatim DVD+R)                    | PB     | 98 ms             | 11 MB/sec      | Stat, RA, OF, WORM |
| Ter  | LTO Streamer<br>(Freecom LTO-920i)           | PB     | 58 sec            | 120 MB/sec     | Stat, SA, OF       |

Pri= Primary, Sec=Secondary, Ter=Tertiary

Vol=Volatile, Stat=Static, Dyn=Dynamic, RA=Random Access, SA=Sequential Access

OL=Online, OF=Offline, WORM=Write Once Read Many



## 8.4 Storage Media

*Detour*

- Alternative to hard-drives: **SSD**
  - Use microchips which retain data in non-volatile **memory** and contain **no moving parts**
- Use the same interface as hard disk drives
  - Easily replacing in most applications possible
- Key components
  - Memory
  - Controller (embedded processor)





## 8.4 SSDs

*Detour*

- The memory:
  - Retains memory even without power
  - Slower than DRAM solutions
  - **Wears down!**
- The controller:
  - **Wear levelling, bad block mapping**, read and write caching, encryption.



## 8.4 SSDs

*Detour*

- **Advantages**
  - Low access time and latency
  - No moving parts → shock resistant
  - Silent
  - Lighter and more energy-efficient than HDDs
- **Disadvantages**
  - Divided into blocks; if one byte is changed the whole block has to be rewritten (write amplification)
  - 10 % of the storage capacity are allocated (spare area)
  - **Limited ability of being rewritten** (between 3000 and 100,000 cycles per cell)
    - Wear leveling algorithms assure that write operations are equally distributed to the cells



## 8.4 SSD based DWs

*Detour*

- Hybrid solutions (OLTP & OLAP) with SSDs
  - Advantages:
    - Up to 10x faster than HDDs
    - Stores persistent data
  - Disadvantages:
    - SSDs wear down: limited ability of being **rewritten**
      - Not really suitable for OLTP
    - Not fast enough



## 8.4 Random Access Memory

# Detour

- Dynamic Random Access Memory(DRAM)
  - Based on the volatile random access memory
    - Sometimes use **internal battery or external power** device to ensure data persistence
  - Ultrafast data access ( $< 10$  microseconds)
  - Rather expensive








## 8.4 DRAM based DWs

*Detour*

- Hybrid solutions (OLTP & OLAP) with **DRAM**
  - **Very high data access speed**
    - Approx. 2000x faster than HDDs and 200x faster than SSDs
  - Once the power supply is gone, the data is lost
    - Use additional batteries and external SSDs to make data persistent
- High cost: Multi-CPU server with **several terabytes of RAM**  
cost 60 000 euros - TU München
  - the HyPer system





## 8.4 HyPer

- Hybrid OLTP & OLAP main memory database system [A. Kemper and T. Neumann, ICDE 2011]
- System goals:
  - Process OLTP transactions at rate of tens of thousands per second, and at the same time
  - Process OLAP queries on up-to-date snapshots of the transactional data





## 8.4 HyPer: OLTP Processing

- Hyper is a main memory database that runs OLTP queries **sequentially**
  - Since all the data is already in the primary memory, the CPU doesn't have to wait for IO operations from the HDD
  - *An order entry or a payment processing* takes only around **10 microseconds**
  - **Tens of thousands/second** of such OLTP transactions can be sustained with such a system




## 8.4 HyPer: OLAP Processing

- Even with such fast hardware, allowing for OLAP to be injected in the OLTP workload would clog the system
  - OLAP queries finish in about 30 milliseconds on such a system
  - Locking the system for 30 milliseconds it blocks – in this time interval **3 000 OLTP queries** could have been executed
- Central idea: **run OLAP on OLTP snapshots**



## 8.4 HyPer: OLAP Processing

- **Snapshot management:** Create a virtual memory snapshot of the OLTP
  - Can be done with OS functionality e.g. *fork* system call for Unix based systems
  - **The fork system call**
    - Copies the complete process state to create a new process
    - The old process is called the **parent** and the new process the **child**
    - Immediately after fork the parent and the child are **identical**






## 8.4 The Fork System Call

- Initially the child's data pages are the same as the parent's pages
  - Except they are made read-only for the child
  - Uses a technique called **copy-on-write**
    - Only data that has been **changed by either parent or child** is physically copied in the child's memory pages
    - The rest of the data is **shadowed through pointers**
    - This way only a fraction of the data is actually transferred!



## 8.4 OLAP Snapshot Management

- The virtual memory snapshot of the OLTP memory pages created by a fork call are used for OLAP
  - The fork process doesn't copy anything!
  - When an OLAP query needs data page *a*, it reads it from the **parent space**






## 8.4 OLAP Snapshot Management

- As soon as  $a$  is **updated** by a transaction from OLTP to  $a'$ , the memory page containing the old data values  $a$  and  $b$  is **copied to the child** (OLAP) space

- The rest of the data remains only in the parent space




- This way we also avoid the problem of **OLAP on changing data**





## 8.4 HyPer: OLAP Processing


- Multiple OLAP queries may make use of multi-core hardware and speed up querying even more
  - Each OLAP query runs as a new fork process
  - Once  $a$  is updated to  $a'$ , the page of  $a$  is copied into the first OLAP query's space
  - Once  $a'$  is updated to  $a''$ , and  $c$  to  $c'$ , the pages of  $a'$  and  $c$  are copied into the second OLAP query's space
  - Etc.





## 8.4 Multi-Thread OLTP Processing

- Multi-core architecture can also speed up OLTP
  - **Reads** can easily be parallelized
  - For inserts, updates and deletes, locks are needed
  - Still most of the workload is represented by **lookups**, table scans and range selects





## 8.4 Real-Time DW Products

*Detour*

- SAP HANA: real-time business intelligence with **in-memory** technology
  - Layer for transactional business applications, planning, data warehousing, and BI tools
  - All enterprise applications OLTP & OLAP share one **column store** database architecture
- Column store?

Get your unfair advantage

Introducing SAP HANA, driving real-time business with in-memory technology

Analyze massive amounts of data

**3600x  
faster\***





# 8.4 SAP HANA


*Detour*

## Row vs. Column Store


(Compressed)

| Document Number | Document Date | Sold-To Party | Order Value | Status | Sales Organization | ... |
|-----------------|---------------|---------------|-------------|--------|--------------------|-----|
| 95769214        | 2009-10-01    | 584           | 10.24       | CLOSED | Germany Frankfurt  | ... |
| 95769215        | 2009-10-01    | 1215          | 124.35      | CLOSED | Germany Berlin     | ... |
| 95779216        | 2009-10-21    | 584           | 47.11       | OPEN   | Germany Berlin     | ... |
| 95779217        | 2009-10-21    | 454           | 21.20       | OPEN   | Germany Frankfurt  | ... |

Row Store



Column Store





## 8.4 SAP HANA

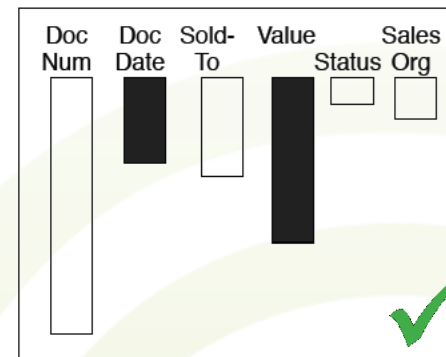
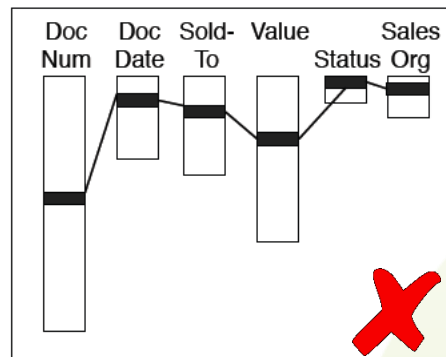
*Detour*

# OLTP vs. OLAP Queries



SELECT \*  
FROM Sales Orders  
WHERE Document Number = '95779216'

SELECT SUM(Order Value)  
FROM Sales Orders  
WHERE Document Date > 2009-01-20

Column Store



Row Store

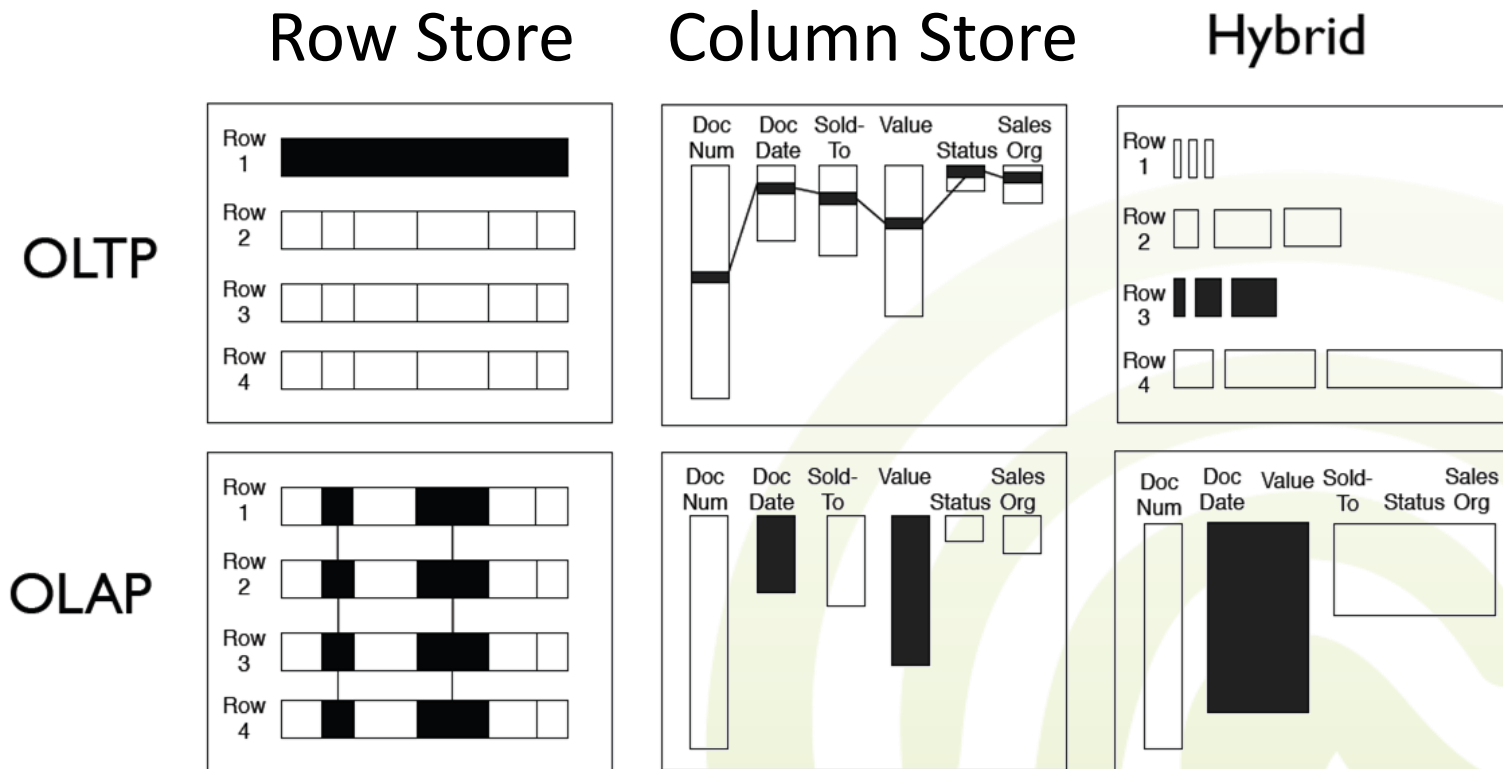




# 8.4 SAP HANA

*Detour*

- HANA: hybrid storage





- How to enable real-time ETL
  - Microbatch and Trickle & Flip for near-time ETL
  - Direct Trickle-Feed with real-time partition – some query contention on the real-time partition
  - External Real-time Data Cache (ERDC) with Just-in-Time data is a pretty complex system
- OLAP and the Changing Data
  - Consistency problems
- OLTP & OLAP - Hybrid Solutions
  - HyPer and SAP HANA





# Next lecture

- Business Intelligence (BI)
  - Principles of Data Mining
  - Association Rule Mining

