

REVIEW OUTLINE

CONTENT

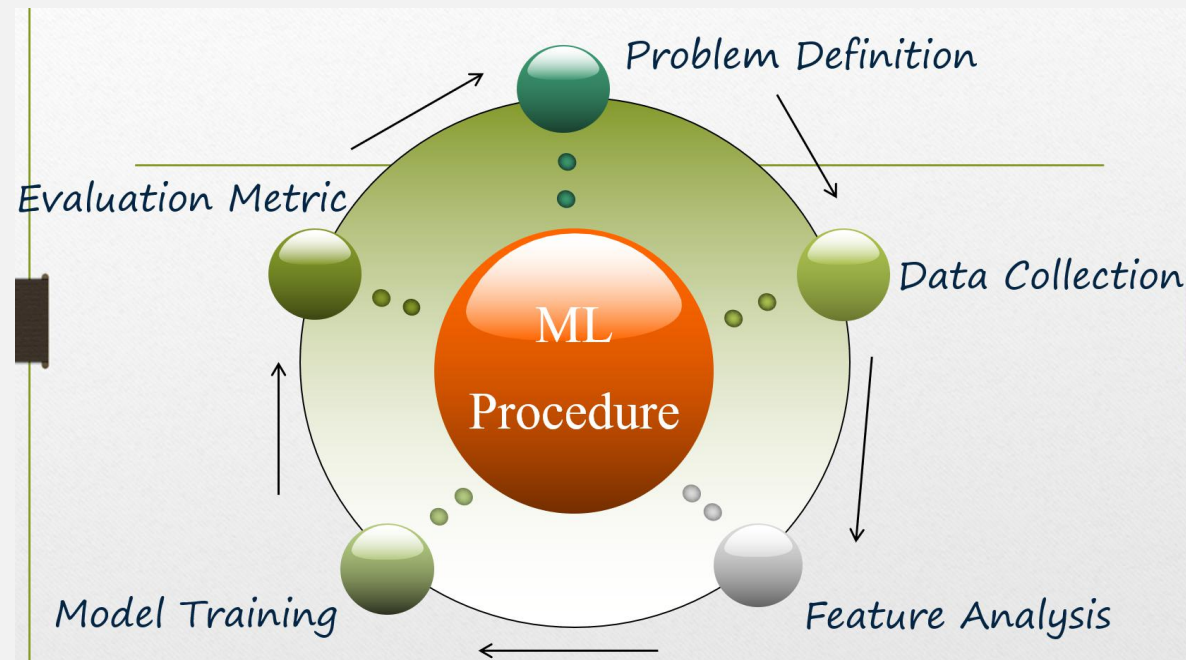
1. Introduction
2. Overview of Supervised Learning
3. Linear Model
4. Decision Tree
5. Neural Network (NN)

INFO OF EXAM

- Closed-book examination: 2024/01/03 (Wednesday) 10:20-12:10
- Presentation (Tencent Meeting, last chance): 2023/12/18 (Monday), 08:30-10:30
- Q & A (Tencent Meeting): 2020/12/25 (Monday), 08:30-10:00
- Problem types (at least 4):
 - Fill-in-the-blank (30%)
 - True/False (10%)
 - Short questions (20%)
 - Comprehensive questions (40%)

I. INTRODUCTION

I.1 General Procedure in ML



I. INTRODUCTION

I.2 Supervised vs. Unsupervised vs. Semi-supervised

- Supervised learning: learn with labeled training data
- Unsupervised learning: learn with unlabeled training data
- Semi-supervised: a small amount of labeled data with a large amount of unlabeled data.
 - Train model with labeled data
 - Use the learned model to predict unlabeled data, then adjust parameter

2. OVERVIEW OF SUPERVISED LEARNING

2.1 Least squares

- Linear model: given a vector of inputs $X^T = (X_1, X_2, \dots, X_p)$

$$\hat{Y} = \hat{\beta}_0 + \sum_{j=1}^p X_j \hat{\beta}_j.$$

bias

weight

- Residual sum of squares (平方残差和):

$$\text{RSS}(\beta) = \sum_{i=1}^N (y_i - x_i^T \beta)^2.$$

$$\text{RSS}(\beta) = (\mathbf{y} - \mathbf{X}\beta)^T (\mathbf{y} - \mathbf{X}\beta),$$

2. OVERVIEW OF SUPERVISED LEARNING

2.2 k-nearest neighbors

- $N_k(x)$ is the neighborhood of x defined by the k closest points x_i

$$\hat{Y}(x) = \frac{1}{k} \sum_{x_i \in N_k(x)} y_i$$

- Common closeness measurement
 - Euclidean distance

$$d(p, q) = \sqrt{(p_1 - q_1)^2 + (p_2 - q_2)^2 + \cdots + (p_i - q_i)^2 + \cdots + (p_n - q_n)^2}.$$

2. OVERVIEW OF SUPERVISED LEARNING

2.3 Loss function vs. Expected prediction error

- Loss function: penalize errors in prediction
 - Squared loss function for regression (f is continuous)
 - Zero-one loss function for classification (f is discrete)
- Expected prediction error : expectation of loss function
- Optimal prediction: to minimize the EPE

2. OVERVIEW OF SUPERVISED LEARNING

2.4 Curse of dimensionality

- When the dimensionality increases, the volume of the space increases so fast that the available data becomes sparse
- The amount of data needed to support the result often grows exponentially with the dimensionality
- Difficult for sampling; local methods inefficient

2. OVERVIEW OF SUPERVISED LEARNING

2.5 Bias–variance decomposition

- **MSE = variance + squared bias**

- Variance: changes in learning performance due to changes in the training set, i.e., impact of data perturbation

$$E_{\mathcal{T}}[\hat{y}_0 - E_{\mathcal{T}}(\hat{y}_0)]^2$$

- Bias: deviation between the expected and real results of the learning algorithm, i.e., fitting ability of learning algorithm

$$E_{\mathcal{T}}(\hat{y}_0) - f(x_0)$$

29

$$\begin{aligned} \text{MSE}(x_0) &= E_{\mathcal{T}}[f(x_0) - \hat{y}_0]^2 \\ &= E_{\mathcal{T}}[\hat{y}_0 - E_{\mathcal{T}}(\hat{y}_0)]^2 + [E_{\mathcal{T}}(\hat{y}_0) - f(x_0)]^2 \\ &= \text{Var}_{\mathcal{T}}(\hat{y}_0) + \text{Bias}^2(\hat{y}_0). \end{aligned}$$

$$\begin{aligned} &E_{\mathcal{T}}[\hat{y}_0 - E_{\mathcal{T}}(\hat{y}_0)]^2 + [E_{\mathcal{T}}(\hat{y}_0) - f(x_0)]^2 \\ &= E_{\mathcal{T}}[\hat{y}_0^2 - 2\hat{y}_0 E_{\mathcal{T}}(\hat{y}_0) + (E_{\mathcal{T}}(\hat{y}_0))^2] + E_{\mathcal{T}}(\hat{y}_0)^2 - 2E_{\mathcal{T}}(\hat{y}_0)f(x_0) + f(x_0)^2 \\ &= E_{\mathcal{T}}(\hat{y}_0^2) - 2E_{\mathcal{T}}(\hat{y}_0)E_{\mathcal{T}}(\hat{y}_0) + E_{\mathcal{T}}(\hat{y}_0)^2 + E_{\mathcal{T}}(\hat{y}_0)^2 - 2E_{\mathcal{T}}(\hat{y}_0)f(x_0) + f(x_0)^2 \\ &= E_{\mathcal{T}}(\hat{y}_0^2) - 2E_{\mathcal{T}}(\hat{y}_0)f(x_0) + f(x_0)^2 \\ &= E_{\mathcal{T}}(\hat{y}_0^2) - 2E_{\mathcal{T}}(\hat{y}_0 f(x_0)) + E_{\mathcal{T}}(f(x_0)^2) \\ &= E_{\mathcal{T}}[(\hat{y}_0^2) - 2\hat{y}_0 f(x_0) + f(x_0)^2] \\ &= E_{\mathcal{T}}[f(x_0) - \hat{y}_0]^2 \end{aligned}$$

3.LINEAR MODEL

3.1 Linear regression

- **Univariate linear regression**

$$f(x_i) = wx_i + b \text{ such that } f(x_i) \approx y_i$$

Where x_i is a scalar

- **Multivariate linear regression**

$$f(\mathbf{x}_i) = \mathbf{w}^T \mathbf{x}_i + b \text{ such that } f(\mathbf{x}_i) \approx y_i$$

Where \mathbf{x}_i is a vector

- **Generalized linear model**


$$y = g^{-1}(\mathbf{w}^T \mathbf{x} + b)$$

Where $g(\cdot)$ is a monotone differentiable function

$$w = \frac{\sum_{i=1}^m y_i (x_i - \bar{x})}{\sum_{i=1}^m x_i^2 - \frac{1}{m} \left(\sum_{i=1}^m x_i \right)^2},$$

$$b = \frac{1}{m} \sum_{i=1}^m (y_i - wx_i)$$

$$\hat{\mathbf{w}}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$



● Posterior probability estimation

$$\ln \frac{y}{1-y} = \mathbf{w}^T \mathbf{x} + b \implies \ln \frac{p(y=1 | \mathbf{x})}{p(y=0 | \mathbf{x})} = \mathbf{w}^T \mathbf{x} + b$$

maximum likelihood method

$$p(y=1 | \mathbf{x}) = \frac{e^{\mathbf{w}^T \mathbf{x} + b}}{1 + e^{\mathbf{w}^T \mathbf{x} + b}},$$

$$p(y=0 | \mathbf{x}) = \frac{1}{1 + e^{\mathbf{w}^T \mathbf{x} + b}}.$$

?

$\prod_{i=1}^m p(y_i | x_i; w, b)$

$$\ell(\boldsymbol{\beta}) = \sum_{i=1}^m \left(-y_i \boldsymbol{\beta}^T \hat{\mathbf{x}}_i + \ln(1 + e^{\boldsymbol{\beta}^T \hat{\mathbf{x}}_i}) \right)$$

Max $\ell(w, b) \leftrightarrow$ Min $\ell(\boldsymbol{\beta})$

$$l(w, b) = \sum_{i=1}^m \ln p(y_i | x_i; w, b)$$

$$= \sum_{i=1}^m \ln [y_i p_1(\hat{x}_i; \boldsymbol{\beta}) + (1 - y_i) p_0(\hat{x}_i; \boldsymbol{\beta})]$$

$$= \sum_{i=1}^m \ln \left[y_i \frac{e^{\boldsymbol{\beta}^T \hat{x}_i}}{1 + e^{\boldsymbol{\beta}^T \hat{x}_i}} + (1 - y_i) \frac{1}{1 + e^{\boldsymbol{\beta}^T \hat{x}_i}} \right]$$

$$= \sum_{i=1}^m \ln \frac{y_i e^{\boldsymbol{\beta}^T \hat{x}_i} + (1 - y_i)}{1 + e^{\boldsymbol{\beta}^T \hat{x}_i}}$$

$$= \begin{cases} y_i = 1, & \sum_{i=1}^m \ln \frac{e^{\boldsymbol{\beta}^T \hat{x}_i}}{1 + e^{\boldsymbol{\beta}^T \hat{x}_i}} = \sum_{i=1}^m (\boldsymbol{\beta}^T \hat{x}_i - \ln(1 + e^{\boldsymbol{\beta}^T \hat{x}_i})) \\ y_i = 0, & \sum_{i=1}^m \ln \frac{1}{1 + e^{\boldsymbol{\beta}^T \hat{x}_i}} = \sum_{i=1}^m (-\ln(1 + e^{\boldsymbol{\beta}^T \hat{x}_i})) \end{cases}$$

3.LINEAR MODEL

3.2 Linear discriminant analysis (LDA)

- Cast the samples onto a straight line
- Project the similar samples as close as possible
- Project the dissimilar samples as far as possible
- For a new sample, determine the class according to the relative position of its projection point.

● Goal: maximize

$$J = \frac{\|w^T \mu_0 - w^T \mu_1\|_2^2}{w^T \Sigma_0 w + w^T \Sigma_1 w}$$

$$= \frac{w^T (\mu_0 - \mu_1) (\mu_0 - \mu_1)^T w}{w^T (\Sigma_0 + \Sigma_1) w}$$



$$J = \frac{w^T S_b w}{w^T S_w w}$$

● Within-class scatter matrix

$$S_w = \Sigma_0 + \Sigma_1$$

$$= \sum_{x \in X_0} (x - \mu_0) (x - \mu_0)^T + \sum_{x \in X_1} (x - \mu_1) (x - \mu_1)^T$$

● Between-class scatter matrix

$$S_b = (\mu_0 - \mu_1) (\mu_0 - \mu_1)^T$$

$$L(w, \lambda) = -w^T S_b w + \lambda (w^T S_w w - 1)$$

$$\frac{\partial L}{\partial w} = -2S_b w + 2\lambda S_w w = 0 \Rightarrow S_b w = \lambda S_w w$$

$$J = \frac{w^T S_b w}{w^T S_w w} \xrightarrow{\text{set } w^T S_w w = 1} \min_w -w^T S_b w \xrightarrow{\text{Lagrange multipliers}} S_b w = \lambda S_w w$$

$w = S_w^{-1} (\mu_0 - \mu_1)$

4.DECISION TREE

4.1 Basic algorithm

1. If all the instances are from exactly one class, then the decision tree is an answer node containing that class name.
2. Otherwise,
 - (a) Define a_{best} to be an attribute with some mechanism
 - (b) For each value $v_{best,i}$ of a_{best} , grow a branch from a_{best} to a decision tree constructed recursively from all those instances with value $v_{best,i}$ of attribute a_{best} .

4.DECISION TREE

4.2 Attribute Selection

- Information gain
- Gain ratio
- Gini index

$$\text{Gain}(D, a) = \text{Ent}(D) - \sum_{v=1}^V \frac{|D^v|}{|D|} \text{Ent}(D^v)$$

$$\text{Gain_ratio}(D, a) = \frac{\text{Gain}(D, a)}{\text{IV}(a)} \quad \text{IV}(a) = - \sum_{v=1}^V \frac{|D^v|}{|D|} \log_2 \frac{|D^v|}{|D|}$$

$$\text{Gini_index}(D, a) = \sum_{v=1}^V \frac{|D^v|}{|D|} \text{Gini}(D^v)$$

4.DECISION TREE

4.3 Bi-partition for continuous value

- Sort n distinct values on a **continuous** attribute a
 $\{a^1, a^2, \dots, a^n\}$
- Split D into D_t^+ and D_t^- w.r.t. a splitting point t
- Candidate set for t

$$T_a = \left\{ \frac{a^i + a^{i+1}}{2} \mid 1 \leq i \leq n - 1 \right\}$$

- Choose the best t $\text{Gain}(D, a) = \max_{t \in T_a} \text{Gain}(D, a, t)$
$$= \max_{t \in T_a} \text{Ent}(D) - \sum_{\lambda \in \{-, +\}} \frac{|D_t^\lambda|}{|D|} \text{Ent}(D_t^\lambda)$$

4.DECISION TREE

4.4 Reweight for missing value

- Information gain:

$$\begin{aligned}\text{Gain}(D, a) &= \rho \times \text{Gain}(\tilde{D}, a) & \text{Ent}(\tilde{D}) &= - \sum_{k=1}^{|\mathcal{Y}|} \tilde{p}_k \log_2 \tilde{p}_k \\ &= \rho \times \left(\text{Ent}(\tilde{D}) - \sum_{v=1}^V \tilde{r}_v \text{Ent}(\tilde{D}^v) \right)\end{aligned}$$

- Reset the *weight* w_x of sample x if need:
 - If x has some value on a , just keep w_x
 - Otherwise, first join x into **each node** corresponding to a^v and then set the weight of x to $\tilde{r}_v \cdot w_x$

4.DECISION TREE

4.5 Random forest

1. For $b = 1$ to B :
 - (a) Draw a bootstrap sample \mathbf{Z}^* of size N from the training data.
 - (b) Grow a random-forest tree T_b to the bootstrapped data, by recursively repeating the following steps for each terminal node of the tree, until the minimum node size n_{min} is reached.
 - i. Select m variables at random from the p variables.
 - ii. Pick the best variable/split-point among the m .
 - iii. Split the node into two daughter nodes.
2. Output the ensemble of trees $\{T_b\}_1^B$.


To make a prediction at a new point x :

Regression: $\hat{f}_{\text{rf}}^B(x) = \frac{1}{B} \sum_{b=1}^B T_b(x)$.

Classification: Let $\hat{C}_b(x)$ be the class prediction of the b th random-forest tree. Then $\hat{C}_{\text{rf}}^B(x) = \text{majority vote } \{\hat{C}_b(x)\}_1^B$.

5. NEURAL NETWORKS (NN)

5.1 Model representation



$$a_1^{(2)} = g(\Theta_{10}^{(1)} x_0 + \Theta_{11}^{(1)} x_1 + \Theta_{12}^{(1)} x_2 + \Theta_{13}^{(1)} x_3)$$

$$a_2^{(2)} = g(\Theta_{20}^{(1)} x_0 + \Theta_{21}^{(1)} x_1 + \Theta_{22}^{(1)} x_2 + \Theta_{23}^{(1)} x_3)$$

$$a_3^{(2)} = g(\Theta_{30}^{(1)} x_0 + \Theta_{31}^{(1)} x_1 + \Theta_{32}^{(1)} x_2 + \Theta_{33}^{(1)} x_3)$$

$$h_{\Theta}(x) = g(\Theta_{10}^{(2)} a_0^{(2)} + \Theta_{11}^{(2)} a_1^{(2)} + \Theta_{12}^{(2)} a_2^{(2)} + \Theta_{13}^{(2)} a_3^{(2)})$$

Vectorized implementation

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{bmatrix} \quad z^{(2)} = \begin{bmatrix} z_1^{(2)} \\ z_2^{(2)} \\ z_3^{(2)} \end{bmatrix}$$

$$z^{(2)} = \Theta^{(1)} x$$

$$a^{(2)} = g(z^{(2)})$$

$$\text{Add } a_0^{(2)} = 1$$

$$z^{(3)} = \Theta^{(2)} a^{(2)}$$

$$h_{\Theta}(x) = a^{(3)} = g(z^{(3)})$$

5. NEURAL NETWORKS (NN)

5.2 Backpropagation algorithm

Backpropagation algorithm

Training set $\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$

Set $\Delta_{ij}^{(l)} = 0$ (for all l, i, j).

For $i = 1$ to m

Set $a^{(1)} = x^{(i)}$

Perform forward propagation to compute $a^{(l)}$ for $l = 2, 3, \dots, L$

Using $y^{(i)}$, compute $\delta^{(L)} = (a_i^{(L)} - y) \cdot g'(z_i^{(L)})$

Compute $\delta^{(L-1)}, \delta^{(L-2)}, \dots, \delta^{(2)}$

reset weights: $\theta_{ji}^l = \theta_{ji}^l - \alpha \cdot \delta_j^{(l+1)} \cdot a_i^{(l)}$

$$a_j^{l+1} = g\left(\sum_{i=1}^{S_l} \theta_{ji}^l \cdot a_i^l\right)$$

$$\delta_i^{(l)} = \frac{\partial E}{\partial z_i^{(l)}} = \sum_j^{N^{(l+1)}} (\delta_j^{(l+1)} \cdot \Theta_{ji}^{(l)} \cdot g'(z_i^{(l)}))$$

5. NEURAL NETWORKS (NN)

5.3 Convolutional Neural Network (CNN)

